

## ER-to-relational Mapping

1/28

So far, have considered mappings for ...

- ER attribute  $\rightarrow$  relational attribute
  - ER entity  $\rightarrow$  relational table
  - ER key  $\rightarrow$  primary key for table
  - n:m relationship  $\rightarrow$  relational table  
(with foreign key for each participating entity plus relationship attributes)
  - 1:n relationship  $\rightarrow$  foreign key plus relationship attributes
  - 1:1 relationship  $\rightarrow$  foreign key plus relationship attributes
- 

## n-way Relationships

2/28

Relationship mappings above assume binary relationship.

If multiple entities are involved:

- n:m generalises naturally to n:m:p:q
  - include foreign key for each participating entity
  - include any other attributes of the relationship
- other multiplicities (e.g. 1:n:m) ...
  - need to be mapped the same as n:m:p:q
  - so not quite an accurate mapping of the ER

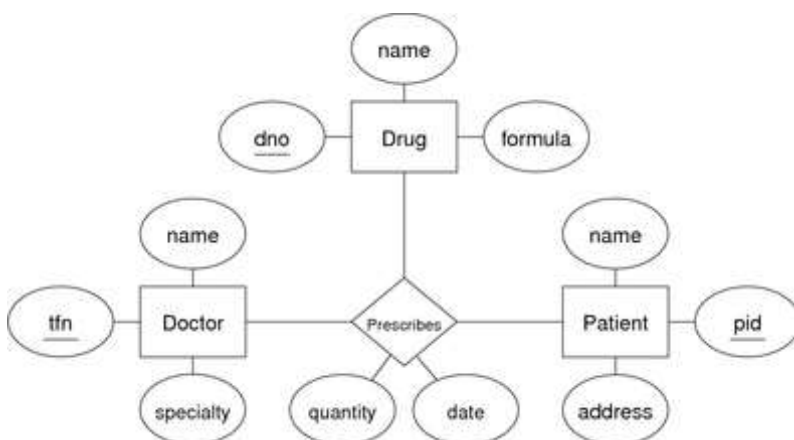
Some people advocate converting n-way relationships into:

- a new entity, and a set of n binary relationships
- 

## Exercise: 3-way relationship

3/28

Translate the following ER design to a relational schema:



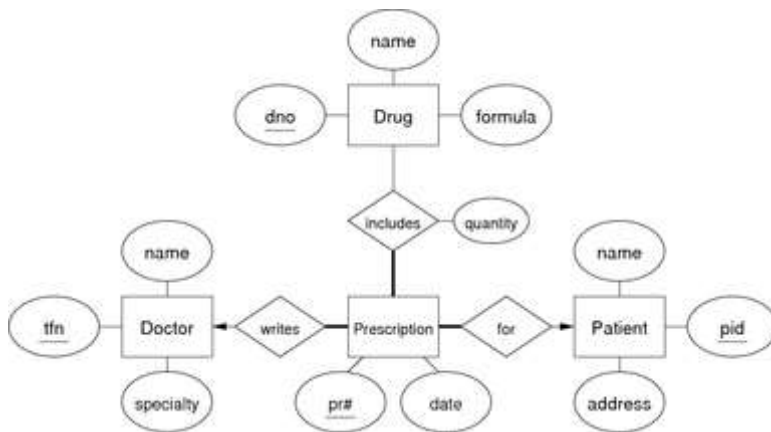
[Solution]

---

## Exercise: Alternative prescription model

4/28

Translate the following ER design to a relational schema:



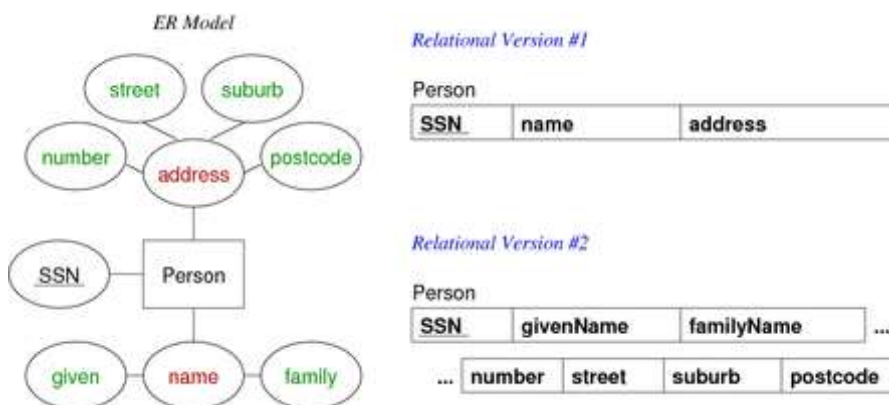
[Solution]

## Mapping Composite Attributes

5/28

Composite attributes are mapped by concatenation or flattening.

Example:

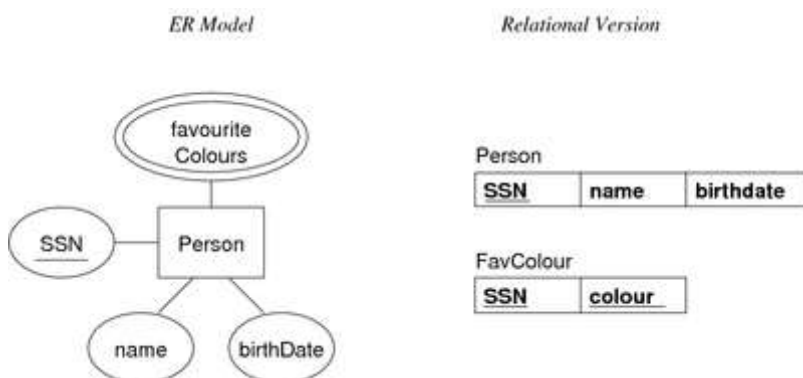


## Mapping Multi-valued Attributes (MVAs)

6/28

MVAs are mapped by a new table linking values to their entity.

Example:



## ... Mapping Multi-valued Attributes (MVAs)

7/28

Example: the two entities

```
Person(12345, John, 12-feb-1990, [red, green, blue])
Person(54321, Jane, 25-dec-1990, [green, purple])
```


would be represented as

```
Person(12345, John, 12-feb-1990)
Person(54321, Jane, 25-dec-1990)
FavColour(12345, red)
FavColour(12345, green)
FavColour(12345, blue)
FavColour(54321, green)
FavColour(54321, purple)
```

## Exercise: Attribute Mappings

8/28


Convert this ER design to relational form:



## ... Mapping Subclasses

11/28

Example of object-oriented mapping:



*Relational Model*

Person		
<b>SSN</b>	<b>name</b>	<b>address</b>

Employee				
<b>SSN</b>	<b>name</b>	<b>address</b>	<b>salary</b>	<b>position</b>

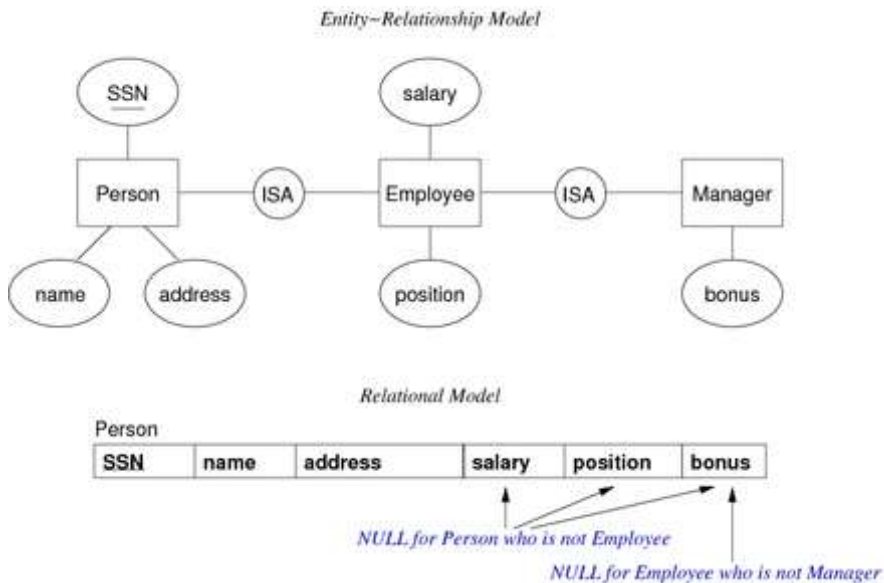
  

Manager					
<b>SSN</b>	<b>name</b>	<b>address</b>	<b>salary</b>	<b>position</b>	<b>bonus</b>

12/28

## ... Mapping Subclasses

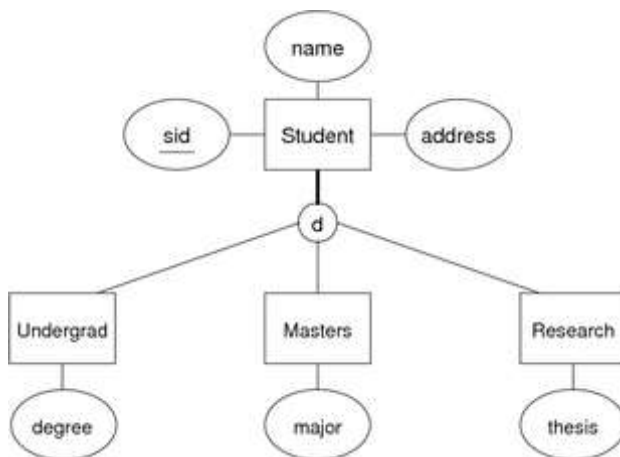
Example of single-table-with-nulls mapping:



## Exercise: Disjoint subclasses

13/28

Translate the following ER design to a relational schema:



Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping

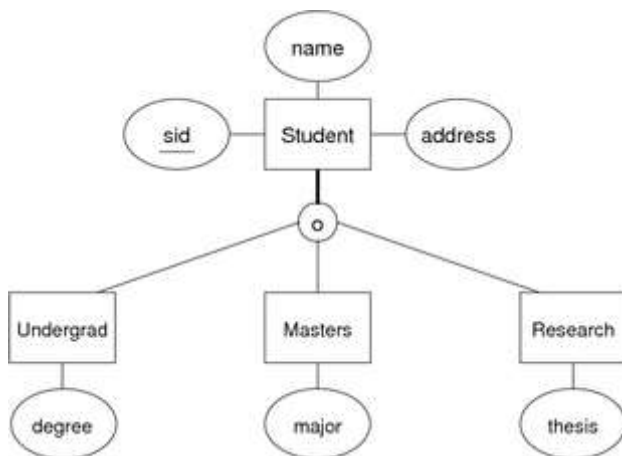
Are there aspects of the ER design that can't be mapped?

[Solution]

## Exercise: Overlapping subclasses

14/28

Translate the following ER design to a relational schema:



Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping

Are there aspects of the ER design that can't be mapped?

[\[Solution\]](#)

## Relational DBMSs

## What is an RDBMS?

16/28

A relational database management system (RDBMS) is

- software designed to support large-scale data-intensive applications
- allowing high-level description of data (tables, constraints)
- with high-level access to the data (relational model, SQL)
- providing efficient storage and retrieval (disk/memory management)
- supporting multiple simultaneous users (privilege, protection)
- doing multiple simultaneous operations (transactions, concurrency)
- maintaining reliable access to the stored data (backup, recovery)

Note: databases provide persistent storage of information

## Describing Data

17/28

RDBMSs implement  $\cong$  the relational model.

Provide facilities to define:

- domains, attributes, tuples, tables
- constraints (domain, key, referential)

Variations from the relational model:

- no strict requirement for tables to have keys
- bag semantics, rather than set semantics
- no standard support for general (multi-table) constraints

## RDBMS Operations

18/28

RDBMSs typically provide at least the following:

- create/remove a database or a schema
- create/remove/alter tables within a schema
- insert/delete/update tuples within a table
- queries on data, define named queries (views)
- transactional behaviour (ACID)


Most also provide mechanisms for

- creating/managing users of the database
- defining/storing procedural code to manipulate data
- implementing complex constraints (triggers)
- defining new data types and operators (less common)

## PostgreSQL Architecture

19/28

PostgreSQL's client-server architecture:



## Using PostgreSQL

20/28

Using your PostgreSQL server in CSE (once installed):

- login to grieg, set up environment, start server
- use psql, etc. to manipulate databases
- stop server, log off grieg

```

wagner$ ssh YOU@grieg
grieg$ priv srvr
grieg$ source /srvr/YOU/env
grieg$ pg start
grieg$ psql mydb
... do stuff with your database ...
grieg$ pg stop
grieg$ exit
  
```

## ... Using PostgreSQL

21/28

PostgreSQL files (helps to understand state of server)

- PostgreSQL home directory ... /srvr/YOU/pgsql903/
- under the home directory ...
  - postgresql.conf ... main configuration file
  - base/ ... subdirectories containing database files
  - postmaster.pid ... process ID of server process

- `.s.PGSQL.5432` ... socket for clients to connect to server
  - `.s.PGSQL.5432.lock` ... lock file for socket
  - PostgreSQL environment settings ... `/srvr/YOU/env`
- 

## Building/Maintaining Databases

---

## Managing Databases

23/28

Shell commands:

- `createdb dbname`
- `dropdb dbname`

(If no dbname supplied, assumes a database called YOU)

SQL statements:

- `CREATE DATABASE dbname`
- `DROP DATABASE dbname`

(Neither of the above is SQL-standard)

---

## ... Managing Databases

24/28

Shell commands (dump/restore):

- `pg_dump dbname > dumpfile`
- `psql dbname -f dumpfile`

(Database dbname is typically created just before restore)

SQL statements (used in dumpfile):

- `CREATE TABLE table ( Attributes+Constraints )`
  - `ALTER TABLE table TableSchemaChanges`
  - `COPY table ( AttributeNames ) FROM STDIN`
- 

## Managing Tables

25/28

SQL statements:

- `ALTER TABLE table TableSchemaChanges`
- `DROP TABLE table(s) [ CASCADE ]`
- `TRUNCATE TABLE table(s) [ CASCADE ]`

(All conform to SQL standard, but all also have extensions)

DROP..CASCADE drops objects which depend on the table

TRUNCATE..CASCADE truncates tables which refer to the table

---

## Managing Tuples

26/28

SQL statements:

- `INSERT INTO` table ( attrs ) `VALUES` tuple(s)
- `DELETE FROM` table `WHERE` condition
- `UPDATE` table `SET` AttrValueChanges `WHERE` condition

AttrValueChanges is a comma-separated list of:

- attrname = expression

Each list element assigns a new value to a given attribute.

---

## Exercise: Generating IDs

27/28

Consider the following schema:

```
create table T (  
    id serial primary key,  
    x integer,  
    y varchar(10)  
);
```

- what does serial actually produce (look in the catalog)?
  - write INSERT statements to add some tuples
  - how could an application program get the generated id?  
(select max(id) from T may not give the correct result; why not?)
- 

## Managing Other DB Objects

28/28

Databases contain objects other than tables and tuples:

- views, functions, sequences, types, indexes, roles, ...

Most have SQL statements for:

- `CREATE` ObjectType name ...
- `DROP` ObjectType name ...

Views and functions also have available:

- `CREATE OR REPLACE` ObjectType name ...

See PostgreSQL documentation Section IV, Chapter I for SQL statement details.

---

Produced: 9 Aug 2016